

KARTIK GOSWAMI

Newcastle, UK | +44 7823605859 | goswamikartik420@gmail.com

GitHub: github.com/kartik703 | Portfolio: kartik703.github.io

PROFILE

AI Engineer passionate about developing robust AI/ML infrastructure and intelligent systems. Experienced in building LLM-powered apps, automating model lifecycles, and optimizing pipelines with MLOps best practices. Adept at working cross-functionally with technologists, data scientists, and stakeholders to deliver scalable, reproducible, and high-impact AI solutions.

TECHNICAL SKILLS

Languages: Python, SQL

Frameworks & Tools: PyTorch, TensorFlow, Hugging Face, MLflow, ONNX, LangChain, LangGraph

MLOps: CI/CD, MLflow, Model Stores, Feature Stores, Docker, FastAPI, GitHub Actions

AI/ML: LLMs, RAG, Transformers, Fine-Tuning, Prompt Engineering, Classification, Forecasting

Compute: GPU Optimization (CUDA-based pipelines), AWS EC2/GPU, Parallel Processing

DevOps/Infrastructure: REST APIs, Vector Stores, Lightweight Containers, Streamlit

EDUCATION

MSc in Data Science and Artificial Intelligence

Newcastle University, UK Sep 2024 - Sep 2025 (Expected)

Focus Areas: ML Systems, MLOps, LLM Infrastructure, NLP and Model Evaluation

PROJECT EXPERIENCE

RAG-Based Market Research Agent

Developed a LangGraph-powered financial analyst bot with document retrieval and grounded response generation.

Stack: LangGraph, OpenAI, FAISS, Python, Streamlit

PDF Summarizer + Continuous Evaluation Loop

LLM summarizer with model performance tracking and prompt versioning for academic + financial

content.

Stack: Hugging Face, T5, MLflow, Prompt Engineering

MLflow Production Pipeline

Full-cycle ML pipeline with data versioning, model training, and real-time scoring with alerts.

Stack: MLflow, Scikit-learn, XGBoost, Docker

AI Job Match Explainer

Built AI-driven pipeline to score CV-job pairs using LLM embeddings and scalable FastAPI interface.

Stack: GPT-4, LangChain, FastAPI

CERTIFICATIONS

Certified Forage and Grassland Professional (British Airways), Mar 2025

IBM: Exploratory Data Analysis for ML (2024)

Google: Data Foundations (2024)

IBM: What is Data Science? (2024)

GUVI: AI for India 2.0

EXTRA-CURRICULAR & IMPACT

- Organizer of MLOps and GenAI engineering bootcamps at Newcastle University
- Open-source contributor to model lifecycle optimization projects (LangGraph + MLflow)
- Delivered 15+ AI systems with reproducibility, testing, and deployment in mind